

## Homework 5: Classification and its Multiclass Extension

There are six questions (30 total points) in this assignment. The minimum increment is 1 point. Please type in your answers directly in the R Markdown file. After completion, **successfully** knitr it as an html file. Submit **both** the html file and the R Markdown file via Canvas. Please name the R Markdown file in the following format: LastName\_FirstName\_HW5.Rmd, e.g. Zhao\_Zifeng\_HW5.Rmd.

### Credit Default Dataset [18 points]

The credit default dataset contains information on ten thousand customers. The aim here is to predict which customers will default on their credit card debt. The data is stored in `Default.csv`. It contains 4 variables, `default`, `student`, `balance` and `income`. We would like to build several statistical models to predict the probability of `default` of a person with given personal information. The data description is as follows.

- `default`: A factor with levels No and Yes indicating whether the customer defaulted on their debt
- `student`: A factor with levels No and Yes indicating whether the customer is a student
- `balance`: The average balance that the customer has remaining on their credit card after making their monthly payment
- `income`: Income of customer

#### Q1 [6 points] Data Partition and Exploration

**Q1(a) [2 points]** Let's correctly read in the data in `Default.csv` and name it as `total_data`.

```
#reading in the total data from the csv file
total_data <- read.csv("/Users/TomTheIntern/Desktop/Mendoza/Mod 2/Advanced Stats/Homework 5/Default.csv")
#setting header as true because the csv has headers
header = TRUE,
#setting an strings to factors
stringsAsFactors = TRUE)
```

**Q1(b) [2 points]** Among the 10000 customers in the dataset, how many of them default?

```
#getting the summary of the data to get the number of defaults
default_count <- summary(total_data$default)

#outputting the results of the data
default_count
```

```
##    No    Yes
## 9667   333
```

Of the 10,000 customers, 333 of them are defaults, or roughly 0.033% of the data.

**Q1(c) [2 points]** Let's partition the data in `total_data` into training (60%) and test data (40%) and store them as R objects `train_data` and `test_data` respectively. Use random seed `set.seed(7)`!

```

#setting the seed to 7 so the process can be repeated
set.seed(7)

#getting the number of rows of the total data
num_rows <- nrow(total_data)

#determining which rows are training rows
training_rows <- sample(1:num_rows, num_rows * 0.60)

#making the training set
train_data <- total_data[training_rows , ]

#making the testing set
test_data <- total_data[-training_rows, ]

```

## Q2 [6 points] Logistic Regression and GAM

**Q2(a) [2 points]** Fit a logistic regression model of default w.r.t. all 3 predictors using the **training data**, name it `lm_full`.

```

#making a logistic regression model
lm_full <- glm(default ~ student + balance + income,
               #setting the data to train data
               data = train_data,
               #setting the family to binomial so that the predictions are scaled correctly
               family = "binomial")

```

**Q2(b) [2 points]** Perform backward selection of `lm_full` via BIC and name the new model `lm_bwd`. Is any variable removed?

```

#making a backward selection model based on the log regression model
lm_bwd <- step(lm_full, direction = "backward")

```

```

## Start:  AIC=1004.23
## default ~ student + balance + income
##
##           Df Deviance    AIC
## - income   1   996.59 1002.6
## <none>      1   996.23 1004.2
## - student  1   998.60 1004.6
## - balance  1  1810.88 1816.9
##
## Step:  AIC=1002.59
## default ~ student + balance
##
##           Df Deviance    AIC
## <none>      1   996.59 1002.6
## - student  1  1007.98 1012.0
## - balance  1  1812.12 1816.1

```

The backward selection process removed the income variable, but kept studentYes and balance.

**Q2(c) [2 points]** Fit a GAM of default w.r.t. all 3 predictors using the **training data**, name it `gam1`. Let's use splines with `df=4` for the numerical predictors, which include `balance` and `income`.

```
#reading in the gam library  
library(gam)
```

```
## Warning: package 'gam' was built under R version 4.4.1
```

```
## Loading required package: splines
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.22-5
```

```
#making a gam model with splines set to balance and income  
gam1 <- gam(default ~ student + s(balance) + s(income),  
             #using the train data  
             data = train_data,  
             #and setting the predictions to binomial  
             family = "binomial")
```

### Q3 [6 points] Model Evaluation (Prediction)

**Q3(a) [2 points]** Use `lm_full` and `gam1` to generate probability predictions for `default` on the test data and store the predicted probability in `lm_full_pred` and `gam1_pred` respectively.

```
#making a series of predictions for the log model using the predict function and the test data  
lm_full_pred <- predict(lm_full, test_data)
```

```
#making a series of predictions for the gam model using the predict function and the test data  
gam1_pred <- predict(gam1, test_data)
```

**Q3(b) [2 points]** Use the `confusionMatrix()` function in the R package `caret` to evaluate the prediction performance of `lm_full` and `gam1`. What are the sensitivity of `lm_full` and `gam1`?

```
#writing in the caret library  
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
#making a confusion matrix based off of the lm_full predictions  
#I need to set the variables to as factor so it avoids being set to a character  
#and then if else the values above 50 to yes and anything lower to no  
confusionMatrix(as.factor(ifelse(lm_full_pred > .50, "Yes", "No")), test_data$default, positive = "Yes")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 3871  87
##           Yes   6  36
##
##           Accuracy : 0.9768
##           95% CI : (0.9716, 0.9812)
##           No Information Rate : 0.9692
##           P-Value [Acc > NIR] : 0.002527
##
##           Kappa : 0.4274
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.29268
##           Specificity : 0.99845
##           Pos Pred Value : 0.85714
##           Neg Pred Value : 0.97802
##           Prevalence : 0.03075
##           Detection Rate : 0.00900
##           Detection Prevalence : 0.01050
##           Balanced Accuracy : 0.64557
##
##           'Positive' Class : Yes
##
```

*#and then doing the same thing with the gam predictions*

```
confusionMatrix(as.factor(ifelse(gam1_pred > .50, "Yes", "No")), test_data$default, positive = "Yes")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 3870  87
##           Yes   7  36
##
##           Accuracy : 0.9765
##           95% CI : (0.9713, 0.981)
##           No Information Rate : 0.9692
##           P-Value [Acc > NIR] : 0.003428
##
##           Kappa : 0.4246
##
## Mcnemar's Test P-Value : 3.693e-16
##
##           Sensitivity : 0.29268
##           Specificity : 0.99819
##           Pos Pred Value : 0.83721
##           Neg Pred Value : 0.97801
##           Prevalence : 0.03075
##           Detection Rate : 0.00900
##           Detection Prevalence : 0.01075
```

```
##          Balanced Accuracy : 0.64544
##
##          'Positive' Class : Yes
##
```

lm\_full has a sensitivity of 0.29268 and an accuracy of 0.9768

gam1 has a sensitivity of 0.29268 and an accuracy of 0.9765

**Q3(c) [2 points]** Note that the sensitivity of `lm_full` and `gam1` are in the range of 30-40%, which means that the models are having a hard time finding the customers that default. This is not surprising considering that most customers do NOT default. One way to improve sensitivity is to use a threshold **lower** than **0.5** to classify the customer. Let's use a new threshold **0.1**, i.e. we think a customer will default if the predicted probability of default  $> 0.1$ .

Use **0.1** as the new classification threshold and calculate the sensitivity for `lm_full` and `gam1`. Did the sensitivity go up? What is the price we need to pay for increasing sensitivity?

```
#altering the confusion matrix to set the standard for a yes prediction to 10%
confusionMatrix(as.factor(ifelse(lm_full_pred > .1, "Yes", "No")), test_data$default, positive = "Yes")
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  No  Yes
##          No 3865  78
##          Yes  12  45
##
##          Accuracy : 0.9775
##          95% CI : (0.9724, 0.9819)
##          No Information Rate : 0.9692
##          P-Value [Acc > NIR] : 0.0009516
##
##          Kappa : 0.4901
##
##          Mcnemar's Test P-Value : 7.303e-12
##
##          Sensitivity : 0.36585
##          Specificity : 0.99690
##          Pos Pred Value : 0.78947
##          Neg Pred Value : 0.98022
##          Prevalence : 0.03075
##          Detection Rate : 0.01125
##          Detection Prevalence : 0.01425
##          Balanced Accuracy : 0.68138
##
##          'Positive' Class : Yes
##
```

```
#and then doing the same thing with the gam predictions
confusionMatrix(as.factor(ifelse(gam1_pred > .1, "Yes", "No")), test_data$default, positive = "Yes")
```

```
## Confusion Matrix and Statistics
##
```

```
##           Reference
## Prediction   No  Yes
##           No 3864  79
##           Yes  13  44
##
##           Accuracy : 0.977
##           95% CI : (0.9719, 0.9814)
##           No Information Rate : 0.9692
##           P-Value [Acc > NIR] : 0.001844
##
##           Kappa : 0.4787
##
## Mcnemar's Test P-Value : 1.229e-11
##
##           Sensitivity : 0.35772
##           Specificity : 0.99665
##           Pos Pred Value : 0.77193
##           Neg Pred Value : 0.97996
##           Prevalence : 0.03075
##           Detection Rate : 0.01100
##           Detection Prevalence : 0.01425
##           Balanced Accuracy : 0.67719
##
##           'Positive' Class : Yes
##
```

The sensitivity shot up for `lm_full` to 0.36585, but the specificity dropped from 0.99845 to 0.99690.

The sensitivity also improved for the gam model to 0.35772, but the specificity dropped from 0.99819 to 0.99665.

## Iris Dataset [12 points]

The famous R.A. Fisher's iris dataset contains the measurements (in centimeters) of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are *Iris setosa*, *versicolor*, and *virginica*.

The data is distributed with R. It contains 5 variables, `Sepal.Length`, `Sepal.Width`, `Petal.Length`, `Petal.Width` and `Species`. We would like to build several statistical models to classify the species of a given iris based on its sepal length and width and petal length and width. The data description is as follows.

- `Sepal.Length`: sepal length
- `Sepal.Width`: sepal width
- `Petal.Length`: petal length
- `Petal.Width`: petal width
- `Species`: species

We first read in the data and partition the data in `total_data` into training (50%) and test data (50%) and store them as R objects `train_data` and `test_data` respectively.

```
library(caret)
total_data <- iris
set.seed(7)
```

```
total_obs <- nrow(total_data)
train_index <- sample(1:total_obs, 0.5*total_obs)
train_data <- total_data[train_index,]
test_data <- total_data[-train_index,]
```

#### Q4 [2 points] Multinomial Logistic Regression

Fit a multinomial logistic regression model of **Species** w.r.t. all 4 predictors using the **training data**, name it **lm1**.

```
#reading in the neural net library
library(nnet)
#making the log regression model
lm1 <- multinom(Species ~ ., data = train_data)
```

```
## # weights: 18 (10 variable)
## initial value 82.395922
## iter 10 value 0.346201
## iter 20 value 0.256072
## iter 30 value 0.132709
## iter 40 value 0.052666
## iter 50 value 0.033212
## iter 60 value 0.021372
## iter 70 value 0.011800
## iter 80 value 0.009370
## iter 90 value 0.009242
## iter 100 value 0.007993
## final value 0.007993
## stopped after 100 iterations
```

#### Q5 [4 points] Multiclass Neural Networks

Fit an NN model of **Species** w.r.t. all 4 predictors using the **training data**, name it **nn1**. For the architecture of NN, let's use one hidden layer with 4 hidden units.

**Q5(a) [2 points]** Let's generate the **training dataset** that are needed for the estimation of NN using the function `model.matrix()` and store it in **x\_train\_nn**. In addition, use the `scale()` function to standardize the predictors by centering with mean and scaling with sd. Let's further combine the dependent variable **Species** with the standardized predictors **x\_train\_nn** generated. Let's further generate the **test dataset** that are needed for the out-of-sample prediction evaluation of NN using the function `model.matrix` and store it in **x\_test\_nn**. Use the `scale()` function to standardize the predictors by centering with mean and scaling with sd.

```
#beginning to make the data into a matrix form
x_train_nn <- model.matrix( ~. , data = train_data)[ , -7]

#pulling the mean of the data from each of the columns
mean <- apply(x_train_nn, 2, FUN = mean)

#pulling the sd from each of the columns
sd <- apply(x_train_nn, 2, FUN = sd)
```

```

#scaling the data on the mean and the sd of the data
x_train_nn <- scale(x_train_nn, center = mean, scale = sd)[ , -1]

#binding label to the data
x_train_nn <- data.frame(Species = as.factor(train_data$Species), x_train_nn)[ , -6]

#renaming the label
colnames(x_train_nn)[1] <- "Species"

#making the test data frame
x_test_nn <- model.matrix(~., data = test_data)[ , -7]

#scaling using the old predictors
x_test_nn <- scale(x_test_nn, center = mean, scale = sd)[ , -1]

#binding label to the data
x_test_nn <- data.frame(Species = as.factor(test_data$Species), x_test_nn)[ , -6]

#renaming the data
colnames(x_test_nn)[1] <- "Species"

```

**Q5(b) [2 points]** Let's fit an NN that has one hidden layer with 4 hidden units and name it `nn1`. Make sure to use random seed `set.seed(7)`!

```

#loading the neural net library
library(neuralnet)

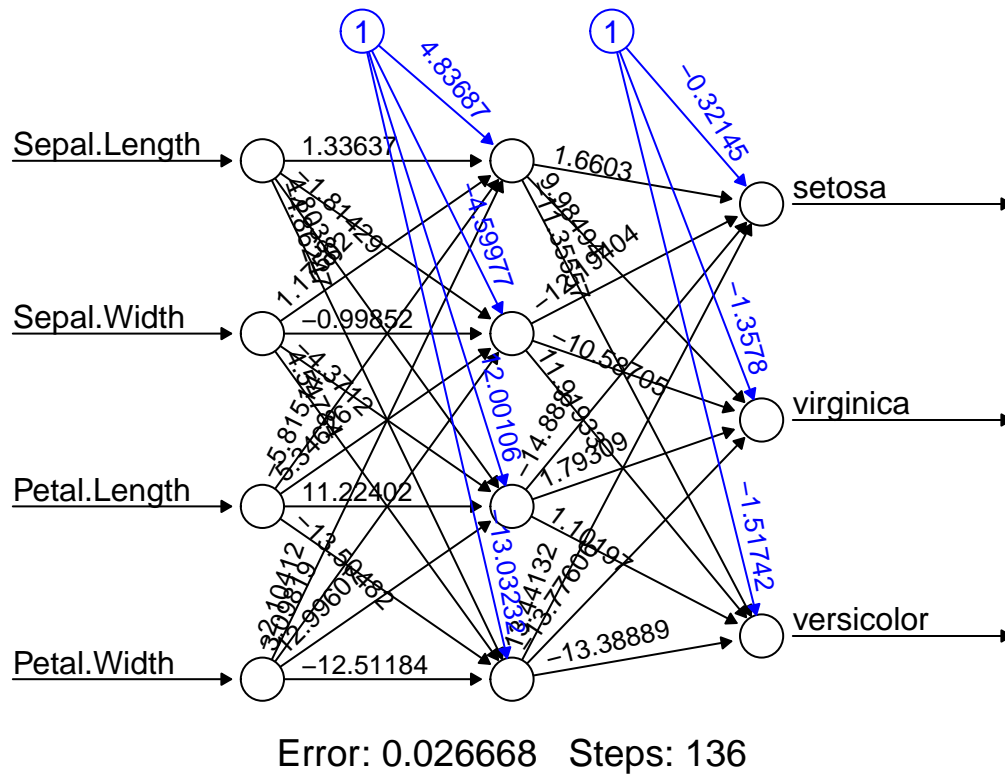
#setting the seed to 7 for repeatability
set.seed(7)

#making the neural net model based on the train data, setting the output to non-linear so it returns mu
nn1 <- neuralnet(Species ~ ., hidden = 4, data = x_train_nn, linear.output = FALSE)

#plotting the neural network model
plot(nn1, rep = "best")

```





#### Q6 [6 points] Model Evaluation (Prediction)

**Q6(a) [2 points]** Use `lm1` and `nn1` to generate probability predictions for `Species` on the test data and store the predicted probability in `lm1_pred` and `nn1_pred` respectively.

```
#making predictions on the log model
lm1_pred <- predict(lm1, newdata = test_data, type = "probs")

#making predictions on the neural network model
nn1_pred <- predict(nn1, newdata = x_test_nn, type = "probs")
```

**Q6(b) [2 points]** Use the `confusionMatrix()` function in the R package `caret` to evaluate the prediction performance of `lm1` and `nn1`.

```
#reading in caret so I can make a confusion matrix
library(caret)

#pulling in the class labels
class_label <- levels(x_train_nn$Species)

#taking the most likely prediction and making it the label
nn1_pred <- factor(class_label[apply(nn1_pred, 1, which.max)])

#making a confusion matrix
confusionMatrix(nn1_pred, x_test_nn$Species)
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction  setosa versicolor virginica
##   setosa      22          0          0
##   versicolor   0          27          0
##   virginica    0           4          22
##
## Overall Statistics
##
##               Accuracy : 0.9467
##               95% CI : (0.869, 0.9853)
##   No Information Rate : 0.4133
##   P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.9196
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: setosa Class: versicolor Class: virginica
## Sensitivity          1.0000          0.8710          1.0000
## Specificity          1.0000          1.0000          0.9245
## Pos Pred Value       1.0000          1.0000          0.8462
## Neg Pred Value       1.0000          0.9167          1.0000
## Prevalence           0.2933          0.4133          0.2933
## Detection Rate       0.2933          0.3600          0.2933
## Detection Prevalence 0.2933          0.3600          0.3467
## Balanced Accuracy    1.0000          0.9355          0.9623
```

```
#taking the most likely preidction and making it the label
lm1_pred <- factor(class_label[apply(lm1_pred, 1, which.max)])

#making a confusion matrix from the data
confusionMatrix(lm1_pred, test_data$Species)
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction  setosa versicolor virginica
##   setosa      22          0          0
##   versicolor   0          27          0
##   virginica    0           4          22
##
## Overall Statistics
##
##               Accuracy : 0.9467
##               95% CI : (0.869, 0.9853)
##   No Information Rate : 0.4133
##   P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.9196
##
```

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.0000           0.8710           1.0000
## Specificity           1.0000           1.0000           0.9245
## Pos Pred Value        1.0000           1.0000           0.8462
## Neg Pred Value        1.0000           0.9167           1.0000
## Prevalence            0.2933           0.4133           0.2933
## Detection Rate        0.2933           0.3600           0.2933
## Detection Prevalence  0.2933           0.3600           0.3467
## Balanced Accuracy      1.0000           0.9355           0.9623
```

**Q6(c) [2 points]** Which statistical model do you prefer, `lm1` or `nn1`? Give reasons.

Answer:

Both models have an accuracy of 0.9467 and perform equally well, so the tiebreaker is picking the model with more of a parsimonious structure, which would be the log model.